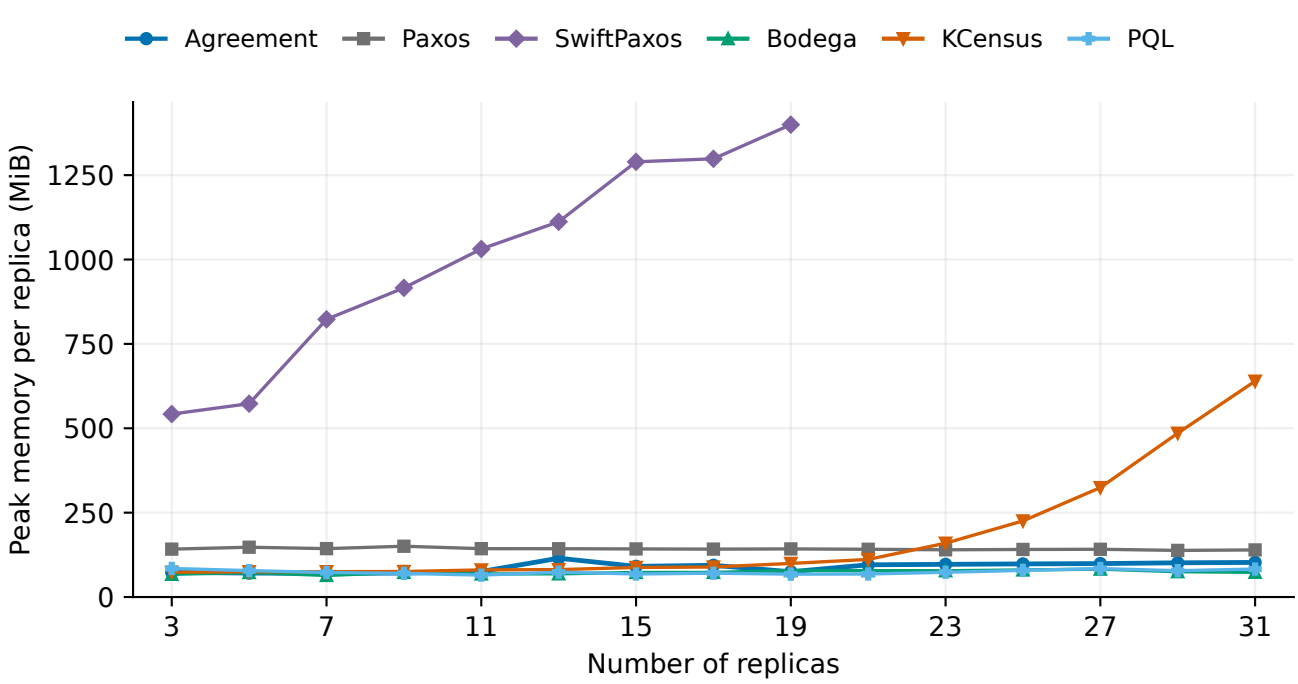

B · AWS-random · 20,000 offered requests/s · complete, fully drained runs only
SwiftPaxos omitted: 21–27 replicas timed out; 29–31 replicas ran out of memory during startup.
Mean of replica lifetime kernel peaks through measurement end; includes startup and warmup.