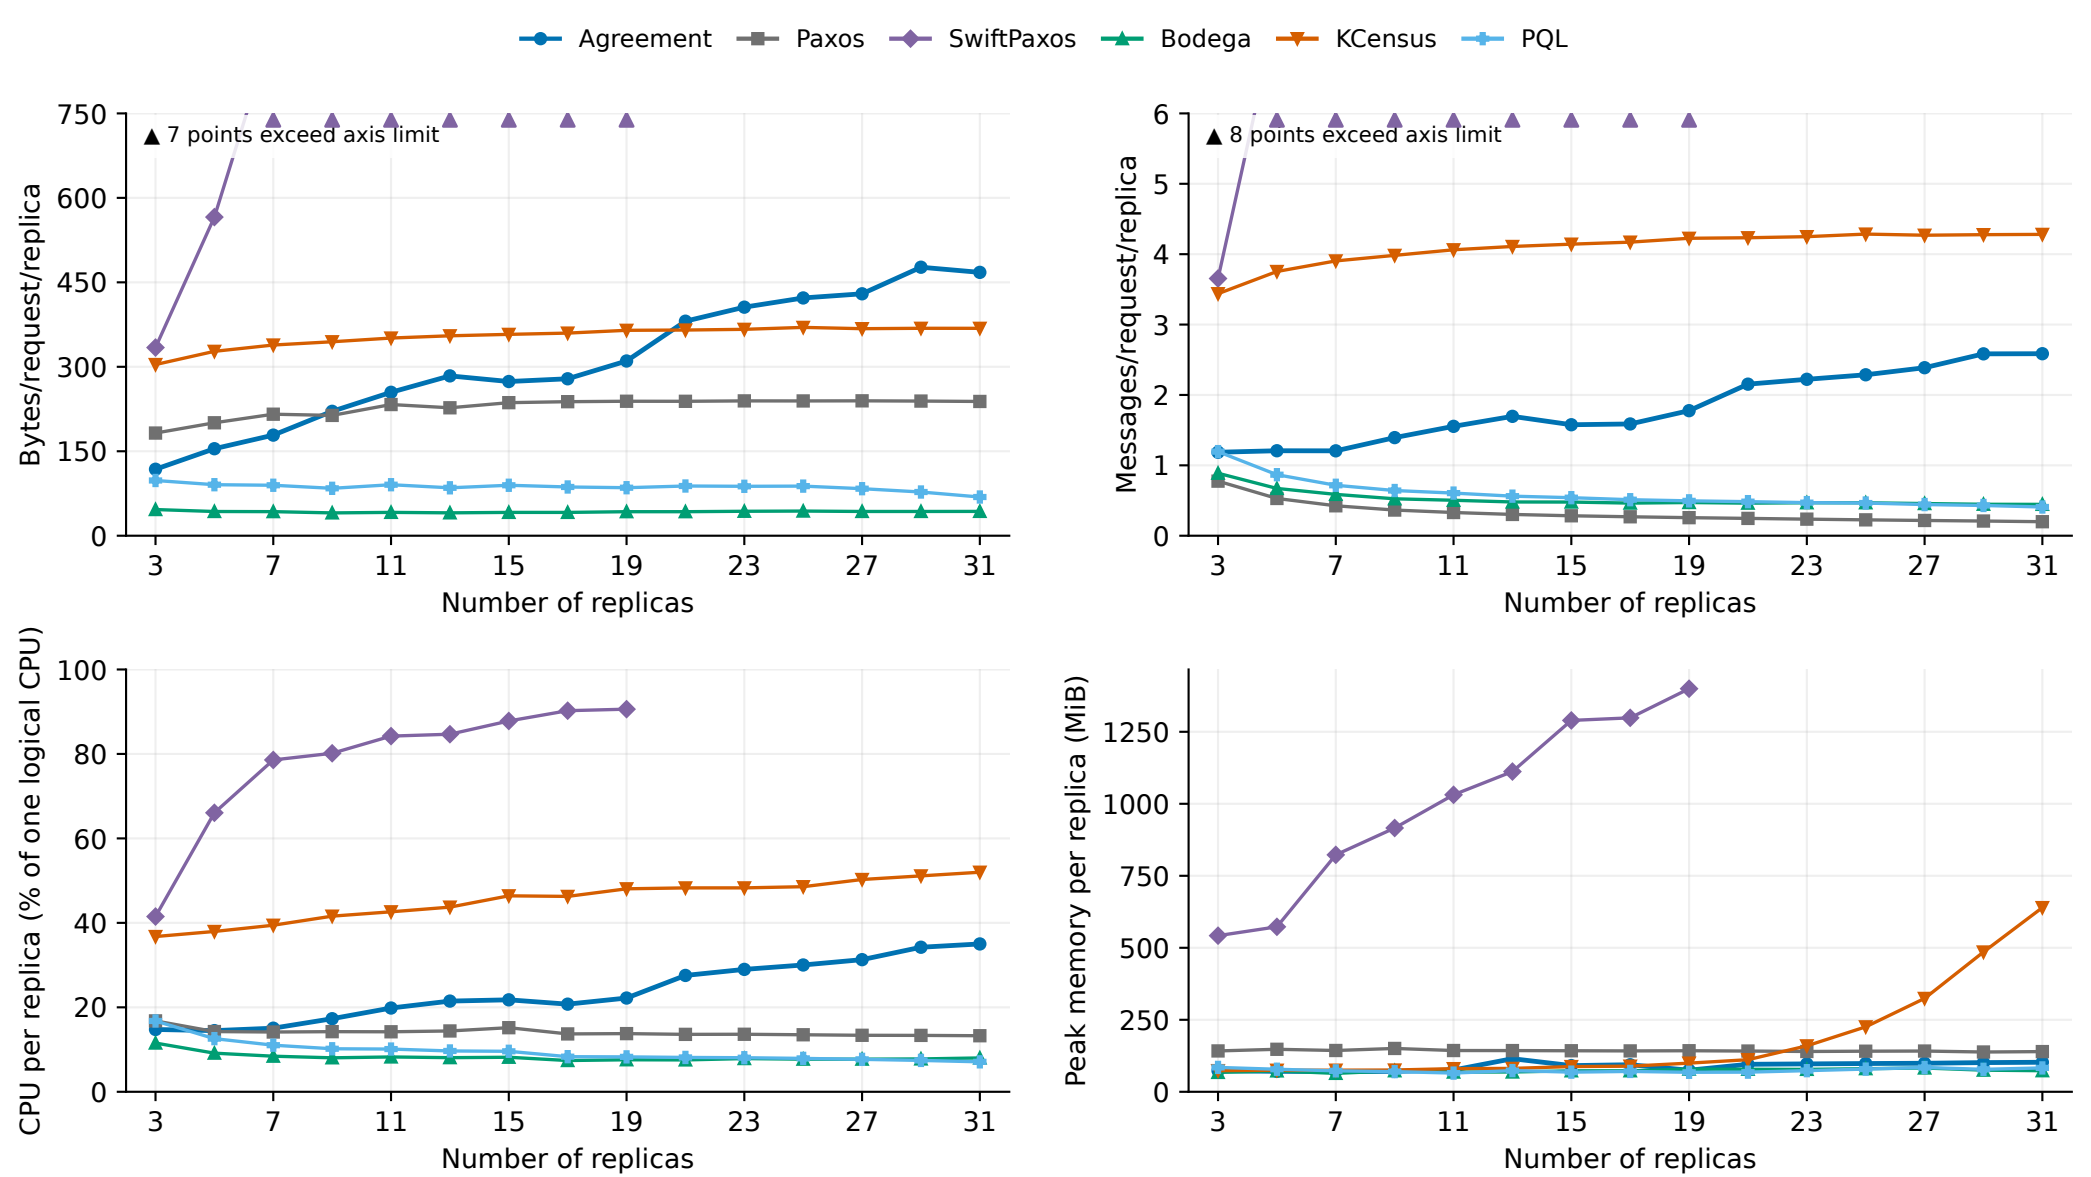

B · AWS-random · 20,000 offered requests/s · complete, fully drained runs only
SwiftPaxos omitted: 21–27 replicas timed out; 29–31 replicas ran out of memory during startup.
Memory: mean of replica kernel peaks, including startup/warmup. ▲ exceeds axis limit; full values remain in CSV.